

lab1

exercise 1

x86-64 属于CISC, 有以下一些特点:

- 指令长度变长
- 寻址方式较多, 比较灵活
- CPU可以不经由寄存器对内存进行操作

而ARM指令集属于RISC,其:

- 指令长度固定
- 寻址方式单一
- 计算基本都在寄存器中完成, 寄存器与内存的通信由load/store来完成
- 寄存器相比前者富余得多, 前者返回地址存储在栈中, 而ARM则有LR寄存器来存储当前函数的返回地址

exercise 2

第一个函数为 `_start()`, 其地址为 `0x00000000000080000`

```
0x00000000000080000 in ?? ()
(gdb) where
#0 0x00000000000080000 in _start ()
Backtrace stopped: not enough registers or memory available to unwind further
```

exercise 3

init 段的地址为 `0x00000000000080000`, 对应上题的 `_start`, 该函数位于 `boot/start.S`,

There are 9 section headers, starting at offset `0x20cf0`:

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
[0]		NULL	0000000000000000	00000000
	0000000000000000	0000000000000000	0 0	0
[1]	init	PROGBITS	0000000000008000	00010000
	000000000000b5b0	0000000000000008	WAX 0 0	4096
[2]	.text	PROGBITS	ffffff000008c000	0001c000
	0000000000001474	0000000000000000	AX 0 0	8
[3]	.rodata	PROGBITS	ffffff0000090000	00020000
	0000000000000118	0000000000000001	AMS 0 0	8
[4]	.bss	NOBITS	ffffff0000090120	00020118
	0000000000008000	0000000000000000	WA 0 0	16
[5]	.comment	PROGBITS	0000000000000000	00020118
	0000000000000032	0000000000000001	MS 0 0	1
[6]	.symtab	SYMTAB	0000000000000000	00020150
	0000000000000858	0000000000000018	7 46	8
[7]	.strtab	STRTAB	0000000000000000	000209a8
	0000000000000308	0000000000000000	0 0	1
[8]	.shstrtab	STRTAB	0000000000000000	00020cb0

```
0000000000000003c 0000000000000000 0 0 1
Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
p (processor specific)
```

```
mrs x8, mpidr_el1
and x8, x8, #0xFF
cbz x8, primary

/* hang all secondary processors before we introduce multi-processors */
secondary_hang:
    bl secondary_hang
```

`mpidr_el1` 记录了 `cpu_id`, 如果其不为0, 则会跳转到 `secondary_hang` 不断循环, 从而实现挂起

exercise 4

```
build/kernel.img:      file format elf64-little

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
  0 init          0000b5b0  0000000000080000 0000000000080000 00010000 2**12
CONTENTS, ALLOC, LOAD, CODE
  1 .text          00001474  ffffffff000008c000 000000000008c000 0001c000 2**3
CONTENTS, ALLOC, LOAD, READONLY, CODE
  2 .rodata        00000118  ffffffff0000090000 0000000000090000 00020000 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
  3 .bss           00008000  ffffffff0000090120 0000000000090120 00020118 2**4
ALLOC
  4 .comment       00000032  0000000000000000 0000000000000000 00020118 2**0
CONTENTS, READONLY
```

除了 `init` 段的VMA和LMA相同外, 其余的section的VMA和LMA都有 `fffffff0000000000` 的offset.

`.comment` 为注释信息, 在可执行文件中应不占空间。

`.init` 中是bootloader的代码, 内存空间还没有初始化, 所以使用LMA寻址。在bootloader准备函数栈和异常向量, 初始化UART, 以及初始化页表并开启MMU后, 进行到位于 `.text` 的kernel部分, 至此开始使用VMA寻址。

exercise 5

```
static int printk_write_num(char **out, long long i, int base, int sign,
                           int width, int flags, int letbase)
{
    char print_buf[PRINT_BUF_LEN];
    char *s;
    int t, neg = 0, pc = 0;
    unsigned long long u = i;

    if (i == 0) {
```

```

    print_buf[0] = '0';
    print_buf[1] = '\0';
    return prints(out, print_buf, width, flags);
}

if (sign && base == 10 && i < 0) {
    neg = 1;
    u = -i;
}
// TODO: fill your code here
// store the digitals in the buffer `print_buf`:
// 1. the last position of this buffer must be '\0'
// 2. the format is only decided by `base` and `letbase` here

//now pointer s should point to the last element of char array `print_buf`
s = print_buf + PRINT_BUF_LEN - 1;
//set the last position to '\0'
*s = '\0';

if (neg) {
    if (width && (flags & PAD_ZERO)) {
        simple_putchar(out, '-');
        ++pc;
        --width;
    } else {
        *--s = '-';
    }
}
//from back to front
while(u){
    --s;
    unsigned int curBit = u % base;
    u = u / base;
    if(curBit <= 9){
        *s = '0' + curBit;
    }else{
        if(letbase)
            *s = 'a' + curBit - 10;
        else
            *s = 'A' + curBit - 10;
    }
}

return pc + prints(out, s, width, flags);
}

```

exercise 6

```

/* Prepare stack pointer and jump to C. */
adr    x0, boot_cpu_stack
add     x0, x0, #0x1000
mov     sp, x0

```

内核栈初始化的函数位于 `boot/start.s`, 其通过 `sp` 指向 `boot_cpu_stack[0]` 的第4096 (0x1000) 字节处为内核栈保留了4096字节的空间

exercise 7

```
0x0000000000080000 in ?? ()
```

```
(gdb) b stack_test
```

```
Breakpoint 1 at 0xffffffff000008c020
```

```
(gdb) c
```

```
Continuing.
```

```
Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
```

```
(gdb) x/10ga $x29
```

```
0xffffffff0000092110 <kernel_stack+8176>: 0x0 0xffffffff000008c018
```

```
0xffffffff0000092120 <kernel_stack+8192>: 0x0 0x0
```

```
0xffffffff0000092130 <kernel_stack+8208>: 0x0 0x0
```

```
0xffffffff0000092140 <kernel_stack+8224>: 0x0 0x0
```

```
0xffffffff0000092150 <kernel_stack+8240>: 0x0 0x0
```

```
(gdb) c
```

```
Continuing.
```

```
Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
```

```
(gdb) x/10ga $x29
```

```
0xffffffff00000920f0 <kernel_stack+8144>: 0xffffffff0000092110 <kernel_stack+8176>
```

```
0xffffffff000008c0d4 <main+72>
```

```
0xffffffff0000092100 <kernel_stack+8160>: 0x0 0xffffffffc0
```

```
0xffffffff0000092110 <kernel_stack+8176>: 0x0 0xffffffff000008c018
```

```
0xffffffff0000092120 <kernel_stack+8192>: 0x0 0x0
```

```
0xffffffff0000092130 <kernel_stack+8208>: 0x0 0x0
```

```
(gdb) c
```

```
Continuing.
```

```
Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
```

```
(gdb) x/10ga $x29
```

```
0xffffffff00000920d0 <kernel_stack+8112>: 0xffffffff00000920f0 <kernel_stack+8144>
```

```
0xffffffff000008c070 <stack_test+80>
```

```
0xffffffff00000920e0 <kernel_stack+8128>: 0x5 0xffffffffc0
```

```
0xffffffff00000920f0 <kernel_stack+8144>: 0xffffffff0000092110 <kernel_stack+8176>
```

```
0xffffffff000008c0d4 <main+72>
```

```
0xffffffff0000092100 <kernel_stack+8160>: 0x0 0xffffffffc0
```

```
0xffffffff0000092110 <kernel_stack+8176>: 0x0 0xffffffff000008c018
```

```
(gdb) c
```

```
Continuing.
```

```
Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
```

```
(gdb) x/10ga $x29
```

```
0xffffffff00000920b0 <kernel_stack+8080>: 0xffffffff00000920d0 <kernel_stack+8112>
```

```
0xffffffff000008c070 <stack_test+80>
```

```
0xffffffff00000920c0 <kernel_stack+8096>: 0x4 0xffffffffc0
```

```
0xffffffff00000920d0 <kernel_stack+8112>: 0xffffffff00000920f0 <kernel_stack+8144>
```

```
0xffffffff000008c070 <stack_test+80>
```

```
0xffffffff00000920e0 <kernel_stack+8128>: 0x5 0xffffffffc0
```

```
0xffffffff00000920f0 <kernel_stack+8144>: 0xffffffff0000092110 <kernel_stack+8176>
```

```
0xffffffff000008c0d4 <main+72>
```

```
(gdb) c
```

```
Continuing.
```

```
Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
```

```
(gdb) x/10ga $x29
```

```

0xffffffff0000092090 <kernel_stack+8048>: 0xffffffff00000920b0 <kernel_stack+8080>
0xffffffff000008c070 <stack_test+80>
0xffffffff00000920a0 <kernel_stack+8064>: 0x3 0xffffffffc0
0xffffffff00000920b0 <kernel_stack+8080>: 0xffffffff00000920d0 <kernel_stack+8112>
0xffffffff000008c070 <stack_test+80>
0xffffffff00000920c0 <kernel_stack+8096>: 0x4 0xffffffffc0
0xffffffff00000920d0 <kernel_stack+8112>: 0xffffffff00000920f0 <kernel_stack+8144>
0xffffffff000008c070 <stack_test+80>
(gdb) c
Continuing.

```

```

Thread 1 hit Breakpoint 1, 0xffffffff000008c020 in stack_test ()
(gdb) x/10ga $x29
0xffffffff0000092070 <kernel_stack+8016>: 0xffffffff0000092090 <kernel_stack+8048>
0xffffffff000008c070 <stack_test+80>
0xffffffff0000092080 <kernel_stack+8032>: 0x2 0xffffffffc0
0xffffffff0000092090 <kernel_stack+8048>: 0xffffffff00000920b0 <kernel_stack+8080>
0xffffffff000008c070 <stack_test+80>
0xffffffff00000920a0 <kernel_stack+8064>: 0x3 0xffffffffc0
0xffffffff00000920b0 <kernel_stack+8080>: 0xffffffff00000920d0 <kernel_stack+8112>
0xffffffff000008c070 <stack_test+80>

```


```


(gdb) x/20x $x29
0xffffffff0000092090 <kernel_stack+8048>: 0xffffffff00000920b0 0xffffffff000008c070
0xffffffff00000920a0 <kernel_stack+8064>: 0x0000000000000003 0x00000000ffffffc0
0xffffffff00000920b0 <kernel_stack+8080>: 0xffffffff00000920d0 0xffffffff000008c070
0xffffffff00000920c0 <kernel_stack+8096>: 0x0000000000000004 0x00000000ffffffc0
0xffffffff00000920d0 <kernel_stack+8112>: 0xffffffff00000920f0 0xffffffff000008c070
0xffffffff00000920e0 <kernel_stack+8128>: 0x0000000000000005 0x00000000ffffffc0
0xffffffff00000920f0 <kernel_stack+8144>: 0xffffffff0000092110 0xffffffff000008c0d4
0xffffffff0000092100 <kernel_stack+8160>: 0x0000000000000000 0x00000000ffffffc0
0xffffffff0000092110 <kernel_stack+8176>: 0x0000000000000000 0xffffffff000008c018
0xffffffff0000092120 <kernel_stack+8192>: 0x0000000000000000 0x0000000000000000

```

FP 处的内存值为 caller function 的 FP 值，FP + 8 处的值为当前函数的 LR 值，FP - 16 处的内存值则为函数参数(没太明白FP - 8 处的值是什么)

exercise 8





SP 的恢复只需将 FP 值重新赋给它即可

exercise 9

```
__attribute__((optimize("O1")))
int stack_backtrace()
{
    printk("Stack backtrace:\n");

    // Your code here.

    //fp of current function(stack_backstrce)
    u64 stackBacktraceFP = read_fp();
    //fp of the caller of current function(stack_backtrace)
    u64 callStackBacktraceFP = *(u64 *)stackBacktraceFP;

    u64 fp = callStackBacktraceFP; // value of fp;

    do{
        printk("LR %lx FP %lx Args ",*(u64 *)(fp + 8),fp);
        u64 argAdrBegin = fp - 16; // address of argList begins at fp - 16
        for(int i = 0; i < 5; ++i){
            printk("%lx ", *(u64 *)(argAdrBegin + 8 * i));
        }
        printk("\n");
        fp = *(u64 *)fp; // value of parent fp
    }while(fp); //value of fp != 0

    return 0;
}
```